Udemy - 1 - Data Engineering - part 2

AWS Data Stores in Machine Learning

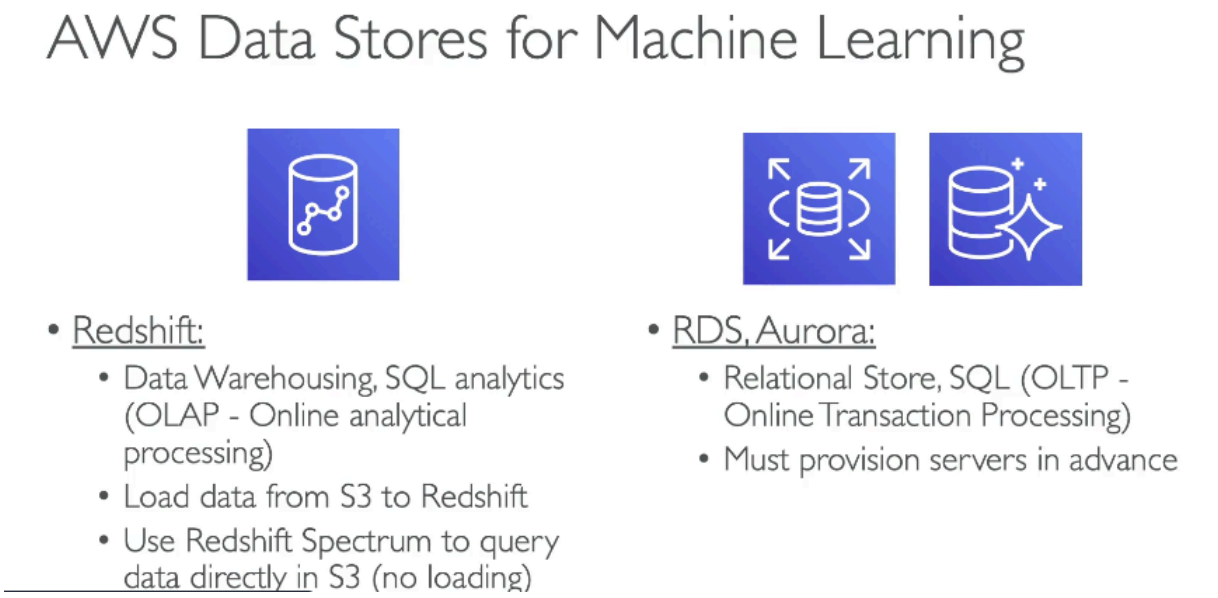

AWS Data Stores for Machine Learning

- ElasticSearch:
  - Indexing of data
  - Search amongst data points
  - Clickstream Analytics

- ElastiCache:
  - Caching mechanism
  - Not really used for Machine Learning

AWS Data Pipelines

It is just an **ORCHESTRATOR**. EC2 instances will be handling the compute

AWS Data Pipeline Features

- Destinations include S3, RDS, DynamoDB, Redshift and EMR
- Manages task dependencies
- Retries and notifies on failures
- Data sources may be on-premises
- Highly available

Why would we use a Data Pipeline?

- Orchestrate and Move data from RDS to S3

Data Pipeline example

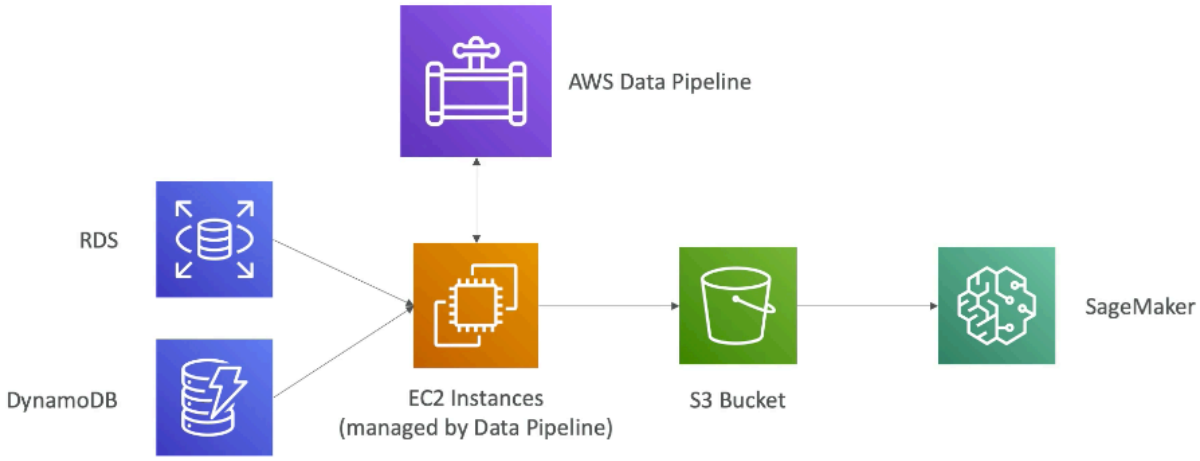

AWS Data Pipeline vs Glue

- Glue:
  - Glue ETL - Run Apache Spark code, Scala or Python based, focus on the ETL
  - Glue ETL - Do not worry about configuring or managing the resources
  - Data Catalog to make the data available to Athena or Redshift Spectrum

- Data Pipeline:
  - Orchestration service
  - More control over the environment, compute resources that run code, & code
  - Allows access to EC2 or EMR instances (creates resources in your own account)

Both are ETL services
- **Glue** is more Apache Spark focused, ETL focused with Transform
- **Data Pipeline** gives us a bit more control, run on EC2 or EMR instances from within our account, gives us a bit more control.

AWS Batch

AWS Batch

- Run batch jobs as Docker images
- Dynamic provisioning of the instances (EC2 & Spot Instances)
- Optimal quantity and type based on volume and requirements
- No need to manage clusters, fully **serverless**
- You just pay for the underlying EC2 instances

- Schedule Batch Jobs using CloudWatch Events
- Orchestrate Batch Jobs using AWS Step Functions

AWS Batch vs Glue

- Glue:
  - Glue ETL - Run Apache Spark code, Scala or Python based, focus on the ETL
  - Glue ETL - Do not worry about configuring or managing the resources
  - Data Catalog to make the data available to Athena or Redshift Spectrum

- Batch:
  - For any computing job regardless of the job (must provide Docker image)
  - Resources are created in your account, managed by Batch
  - For any non-ETL related work, Batch is probably better

DMS - Database Migration Service

DMS – Database Migration Service

- Quickly and securely migrate databases to AWS, resilient, self healing
- The source database remains available during the migration
- Supports:
  - Homogeneous migrations: ex Oracle to Oracle
  - Heterogeneous migrations: ex Microsoft SQL Server to Aurora
- Continuous Data Replication using CDC
- You must create an EC2 instance to perform the replication tasks

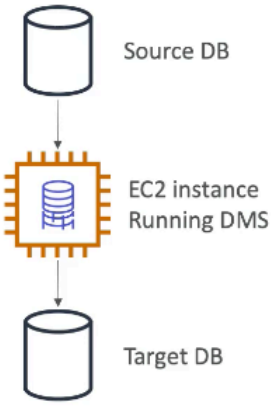

AWS DMS vs Glue

- Glue:
  - Glue ETL - Run Apache Spark code, Scala or Python based, focus on the ETL
  - Glue ETL - Do not worry about configuring or managing the resources
  - Data Catalog to make the data available to Athena or Redshift Spectrum

- AWS DMS:
  - Continuous Data Replication
  - No data transformation
  - Once the data is in AWS, you can use Glue to transform it

AWS Step Functions

Define/design workflow

AWS Step Functions

- Use to design workflows
- Easy visualizations
- Advanced Error Handling and Retry mechanism outside the code
- Audit of the history of workflows
- Ability to "Wait" for an arbitrary amount of time
- Max execution time of a State Machine is 1 year

Example:

Step Functions – Examples
Train a Machine Learning Model

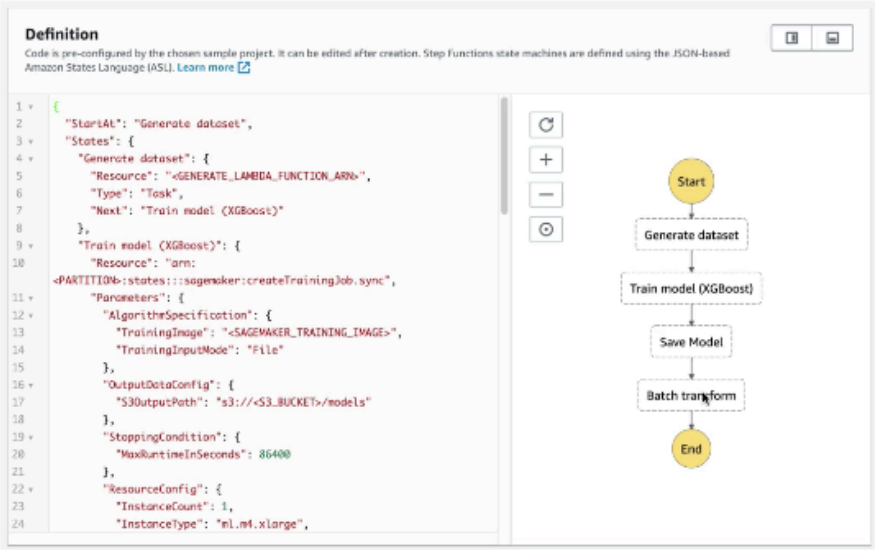

Step Functions – Examples
Tune a Machine Learning Model

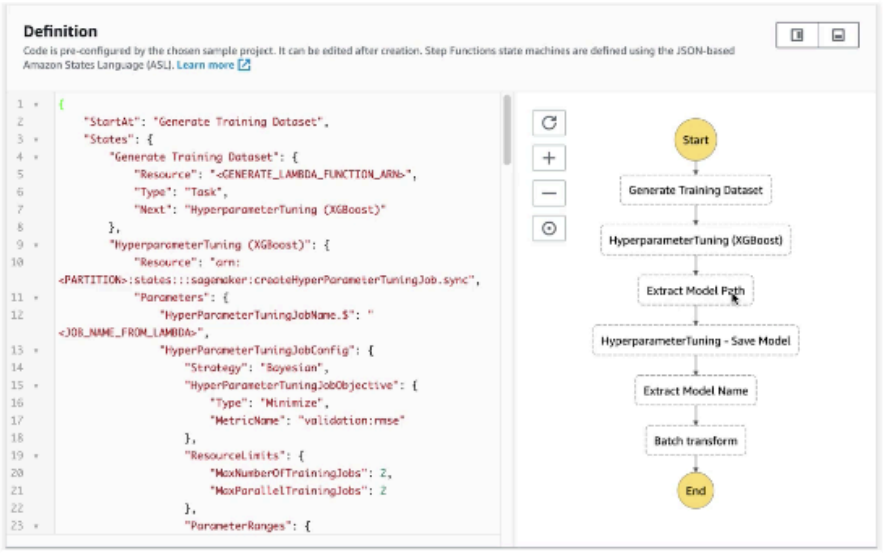

Step Functions – Examples
Manage a Batch Job

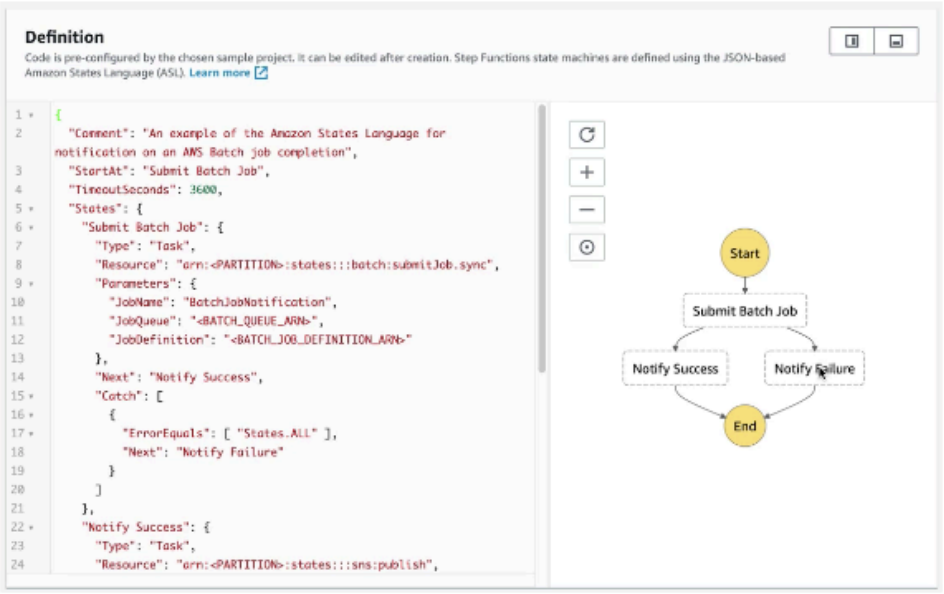

Summarizing All these services - Data Engineering Pipelines

Full Data Engineering Pipeline
Real-Time Layer

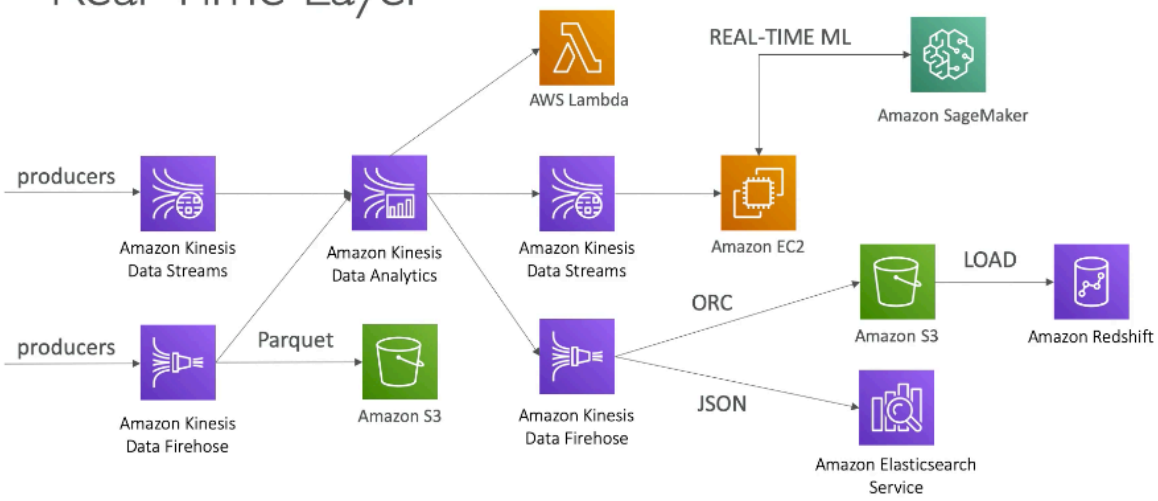

Full Data Engineering Pipeline
Video Layer

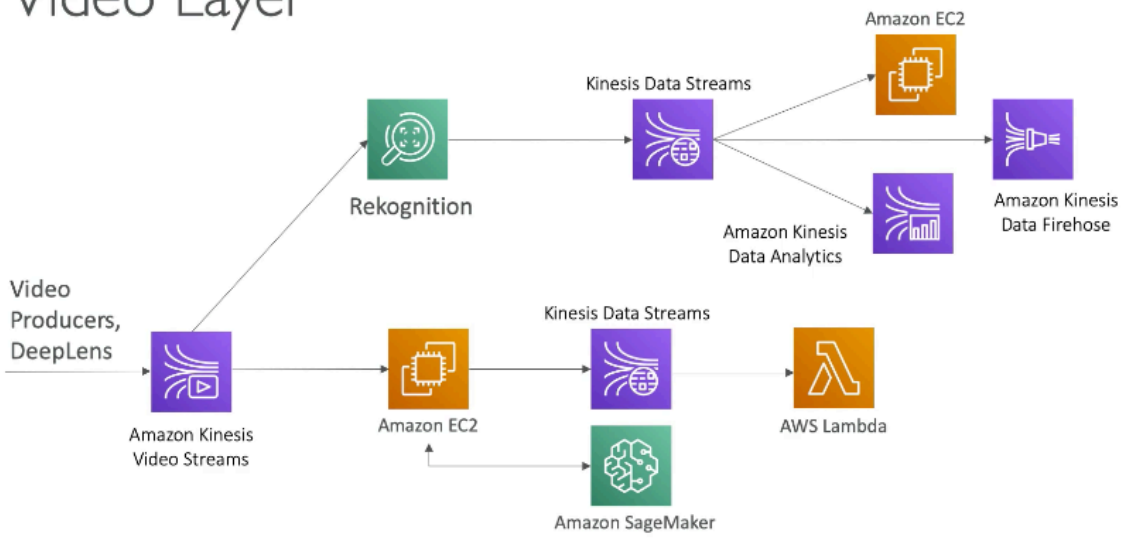

Full Data Engineering Pipeline
Batch Layer

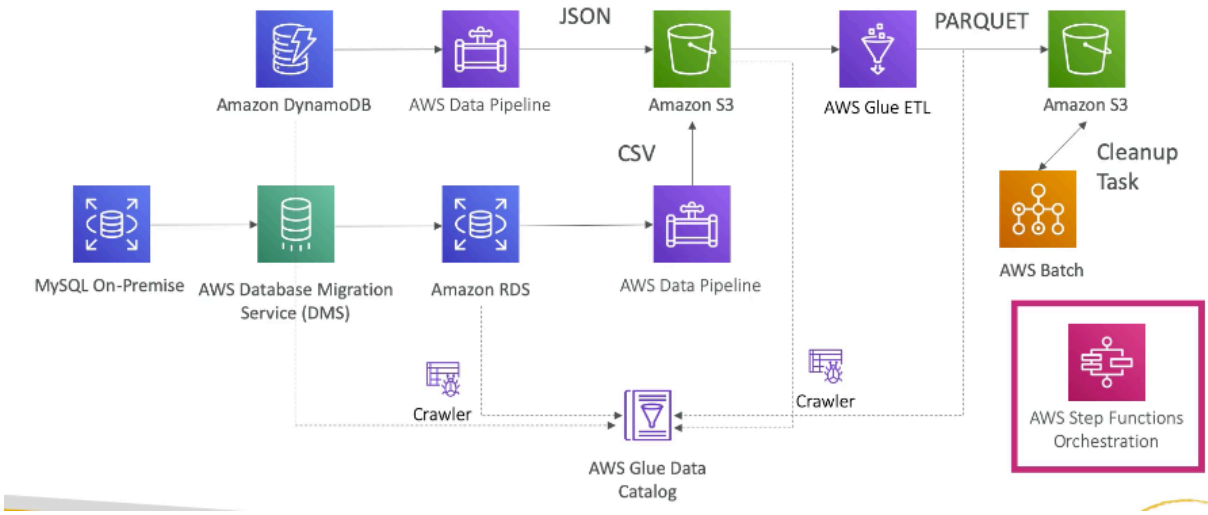

Full Data Engineering Pipeline
Analytics layer

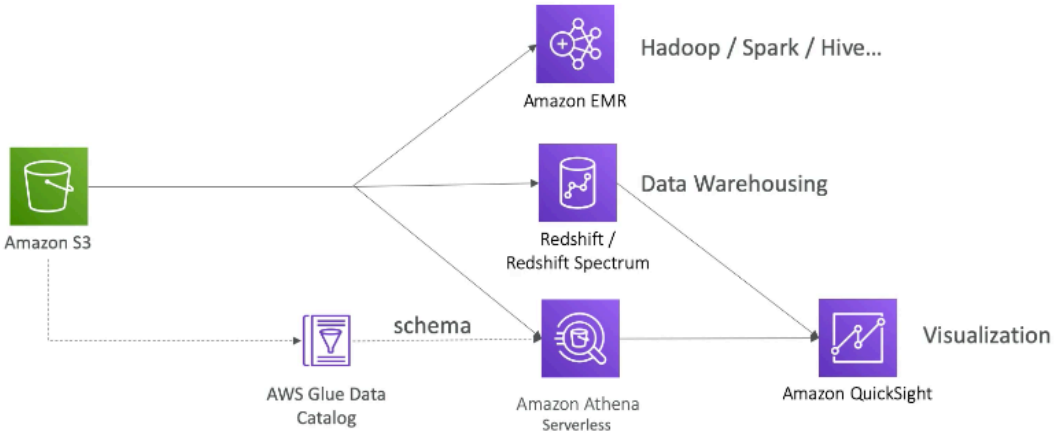

Data Engineering Summary

Here's a quick summary of all the services we've mentioned

- Amazon S3: Object Storage for your data
- VPC Endpoint Gateway: Privately access your S3 bucket without going through the public internet
- Kinesis Data Streams: real-time data streams, need capacity planning, real-time applications
- Kinesis Data Firehose: near real-time data ingestion to S3, Redshift, ElasticSearch, Splunk
- Kinesis Data Analytics: SQL transformations on streaming data

Quiz - Data Engineering

Question 1:

What is the simplest way to manage automating the archiving or deletion of old data in your S3 data lake?

- ○ Write a script that runs periodically using the boto3 API
- ◉ Use S3 Lifecycle Rules
- ○ Use S3 bucket policies
- ○ Use an S3 partitioning strategy

Question 2:

A Kinesis Data Stream's capacity is provisioned by *shards*. What is the maximum throughput of a single shard?

- ○ 100MB / s or 100 messages / s
- ○ 100 MB / s or 1000 messages / s
- ● 1 MB / s or 1000 messages / s
- ○ 1000 MB / s or 100 messages / s

Question 3:

Which Amazon service is appropriate for connecting video data from cameras to backend systems to analyze that data in real time?

- ○ Rekognition
- ○ SageMaker
- ● Kinesis Video Streams
- ○ DeepLens

Question 4:

What is the underlying platform for Glue ETL?

- (●) A serverless Apache Spark platform
- ( ) Amazon Redshift
- ( ) Amazon RDS
- ( ) SageMaker

Question 5:

Which AWS data store provides a highly scalable data warehouse (for OLAP) that can query your S3 data lake directly?

- ( ) Amazon RDS
- (●) Amazon Redshift
- ( ) DynamoDB
- ( ) Elasticsearch

When using Redshift Spectrum, Redshift can query S3 data directly - in addition to many other data sources.